

Big picture (from findings.txt)

- Rows: 90 total. Failures: 18 rows (20%) are failures, defined as error or empty OCR text —in your current CSV, they're all empty OCR rather than explicit errors.
- Processing time (66 rows have a time value): median 8.40 s, mean 23.66 s, p90 39.58 s, max 392.24 s → most images are fast, but there are a few very slow outliers.
- Longest OCR texts: 285, 263, 220 characters.
- Very short outputs: 18 rows have ≤ 3 characters; empty rows are also 18 (same set).
- Relevance to Vietnamese: 75.6% ASCII-only → strong sign that diacritics are often dropped by OCR.
- Text length vs time: $r = 0.429$ ($n=66$) → a moderate positive relationship: longer text tends to take longer, but outliers also matter.

FYI: the script auto-picked detected_text_salvaged as text and unnamed: 5 as your time column; 66/90 rows had non-null time. That's why some time plots are present and some aren't.

What each plot tells you

1) Detected text length (histogram)

- The mass is around ~40–60 characters, with a long tail to 200–285 characters. This means most pages yield short-ish snippets, but the system can capture long passages on some pages. This matches the “longest texts” numbers above.

2) Processing time (histogram)

- Most items finish under ~40 s (p90 ≈ 39.6 s), but a few outliers spike up to ~392 s. These rare long jobs dominate the mean (23.7 s) vs median (8.4 s). In practice, median is a better “typical” latency; the tails are what you'll want to triage.

3) Text length vs processing time (scatter)

- The cloud is low and left (short text, short time), with points rising as text gets longer. That's your $r = 0.429$ correlation showing up. The several very long time points are likely specific tough images (e.g., heavy noise/blur, dense text regions) rather than length alone.

What this means for OCR pipeline

- Primary failure mode: Empty OCR (20%). Investigate image quality gates (blur/contrast/noise), retry strategies (alternate binarization/thresholding), and whether any pre-crop yields blank regions.
- Latency profile: Typical experience is fast (median ~8.4 s), but a small number of cases are very slow; these deserve targeted debugging/alerts.
- Language fidelity: High ASCII-only share (75.6%) suggests lost Vietnamese diacritics—consider language-aware OCR models, Vietnamese-focused post-processing, or better page binarization for faint accents.
- What we can't break down (yet): Your CSV has no folder/result/crops columns—so we can't show failure by scenario, top error strings, or crops↔time effects. (The detector recorded folder_col=None, result_col=None, crops_col=None.)